

# Team Buddie

Topic: Reproducing a Paper (Casual Topic Mining)

*Katie Shin, Sandeep Venkata, Aman Gupta*

# General Idea

- Can we combine probabilistic topic models and causal analysis with external time series to find topics in a corpus of documents that are both coherent semantically and correlated with the time series?
- Reproduced this paper by utilizing two data sets from the paper:
  - Corpus: NYT Articles
  - Time Series: Iowa Presidential Stock Markets
  - Looking for topics that specifically caused support for Bush or Gore to change.

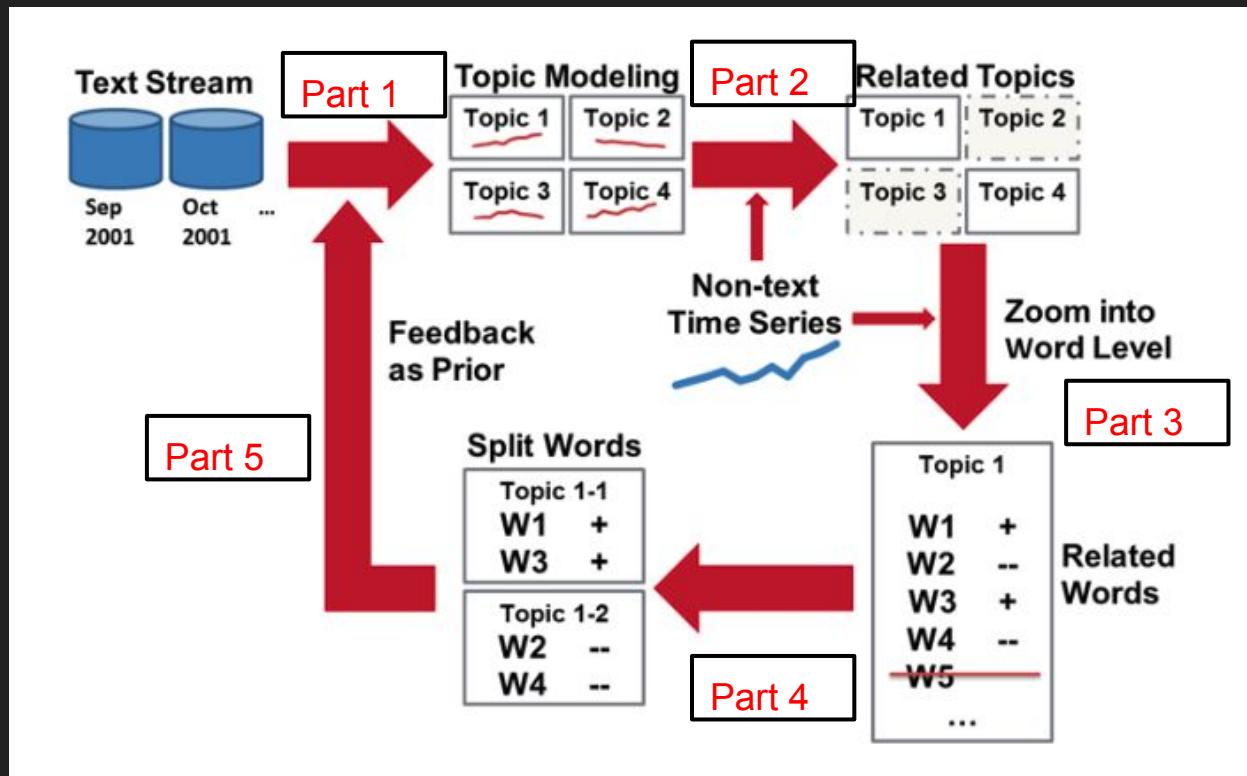
Based on: Hyun Duk Kim, Malu Castellanos, Meichun Hsu, ChengXiang Zhai, Thomas Rietz, and Daniel Diermeier. 2013. Mining causal topics in text data: iterative topic modeling with time series feedback. In Proceedings of the 22nd ACM international conference on Information & Knowledge Management (CIKM '13). Association for Computing Machinery, New York, NY, USA, 885–890. DOI:<https://doi.org/10.1145/2505515.2505612>

# How to Use Software


- Pull from github repo <https://github.com/97agupta/CourseProject>
- Install the necessary libraries, with the most notable one being
  - pip install plsa
- Place NYT corpus in the correct folder format
  - Can be retrieved from <https://catalog ldc.upenn.edu/LDC2008T19>
  - Folder: CourseProject/data/<month>/<day>/<filename>.xml
- To run the iterative code, run 'python3 main.py'
- To run sections of the iterative code, run 'python3 <filename>.py'

(Note: This takes hours to run, and therefore we have provided a sample run w/ end-results here)

# Implementation

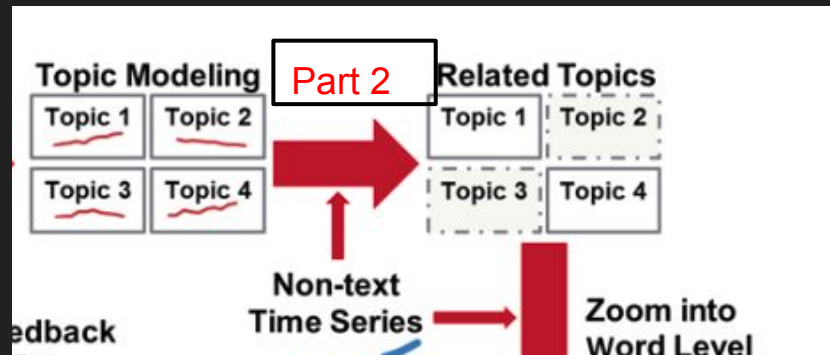


# Implementation



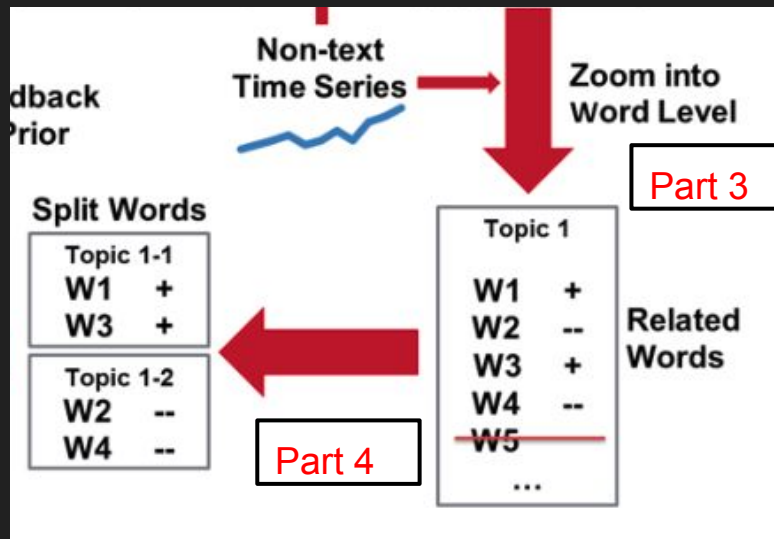
- For each group of xml files that belong to a single day (05/01/2000 - 10/31/2000), we ran and picked the best model out of 5 PLSA models.
  - From the top model, we extract the top 5 highest probability words for each date that will be considered the topic.
- Library used: [PLSA](#)

# Implementation (Part 2)



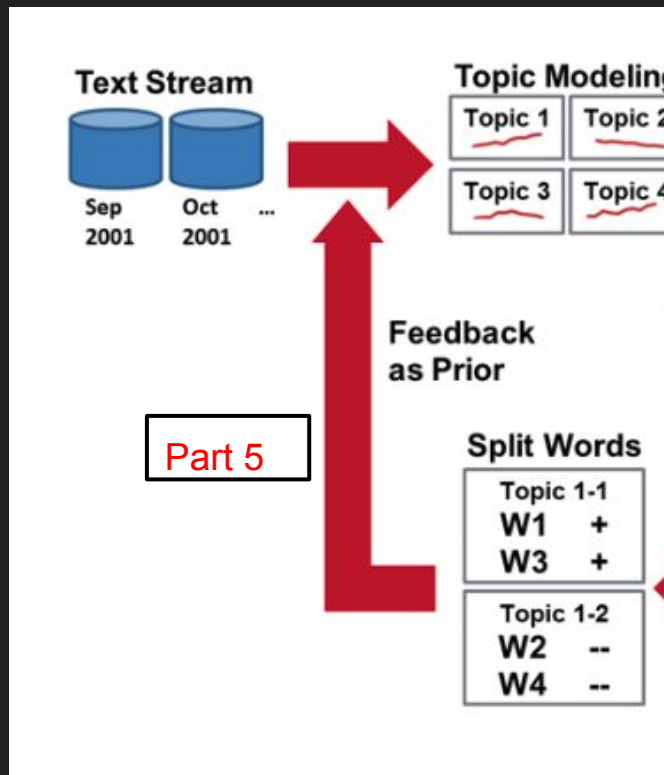
- Based on topics identified during PLSA, we identify related topics by running a Granger Test between the change in probability of topics and the normalized price of one candidate.
  - We use a lag of 5 days & create stationary time series for both of our Granger Test Variables.
  - Not every topic has enough data points for testing, so we skip over topics where the Granger Test fails
  - Based on the highest F-value for the each Granger Test, we determine if a topic is or isn't correlated with the external time series.

# Implementation (Parts 3 & 4)



- From the best PLSA model (refer to Part1), we retrieved the top 20 relevant words (i.e. highest probability) per topic that were deemed related from the Granger test.
- For each word series, we run a pearson coefficient test comparing the external time series and the frequency of top words of each topic.
  - We then segment these words into negative and positive correlations, returning words that when combined, meet our probability threshold (0.75)
  - These words and then used as a Prior for the PLSA algorithm.

# Implementation (Part 5)



- We used the [PyPI](#) package for PLSA but it didn't offer support for including priors to guide the PLSA as per user inputs.
- So we enlisted the [source code](#) and extended the PLSA algorithm to overwrite the M step.
- The library essentially uses numpy 'einsum' to efficiently perform multi-dimensional matrix operations.
- We modified the M-step of the algorithm to include the parameters ( $\mu$ ) and the pseudo counts.
- The key part was to look for only the words that occurred in the day.



# Results (1st Iteration) PLSA without prior

- After running our first PLSA on the corpus of NYT documents, we found 343 topics. Of these topics, only 4 topics showed causality with our external time series: bush, gore, campaign, and clinton.
- For each of these topics, we extracted the top words and segmented each topic into two topics made up of positively and negatively influenced words. These new topics were used as a prior for our second iteration. Here is an example of the Gore topic, being split into two:

gore \_positive

	Word	Probability
0	debate	0.4671585084138510
1	oil	0.24000592562023600
2	bush	0.02037089961487880
3	teacher	0.018116499244389200

gore \_negative

	Word	Probability
0	convention	-0.2445556313924940
1	party	-0.19336969916928100
2	delegate	-0.18026760013699100
3	missile	-0.07297529159633
4	stock	-0.0063307171624487500

# Results (2nd Iteration)

- With a new iteration using PLSA with prior, we find the following relevant topics using our granger test:

- Teacher, oil, drug, debate

```
for key in relevent_topic:  
    print(key)
```

```
teacher  
drug  
oil  
debate
```

- This is a clear improvement over our original PLSA, as it delves deeper into the issues that moved the campaign likelihoods for each candidate

- Bush, Gore, Campaign, Clinton

-- This clearly shows that our model is finding topics that are more related to the campaign issues & with further iterations we expect this to further improve.

# Next Steps

- We would want to run further iterations on these topics.
- Additionally, instead of using the full NYT corpus we could identify paragraphs that mention Bush, Gore, or Presidential to create a smaller corpus from which to topic mine.
- We would want to find a more robust method of comparing and understanding the different F-Scores from each lag of our Granger Test.

Thanks!



(This is Buddie)